

Analysis of Algorithms 1 (Fall 2011) Istanbul Technical University Computer Eng. Dept.



Chapter 13 Red-Black Trees

Last updated: December 03, 2009

Purpose

Review Binary Search Trees

Introduce Red-Black Trees

Review 2-3 and 2-3-4 trees

Rotations and other operations in RB
Trees

Outline

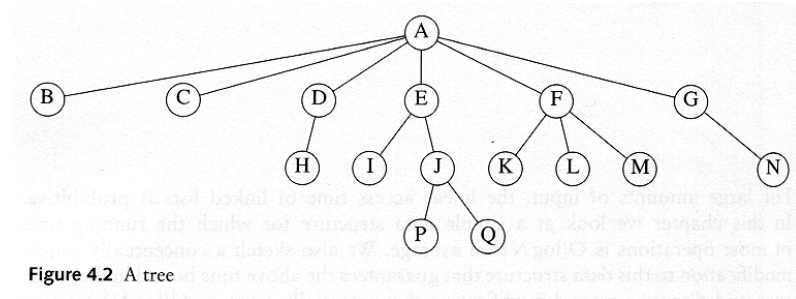
Binary Search Tree (BST) review

Red and Black Trees

2-3 and 2-3-4 trees

Operations on Red and Black Trees

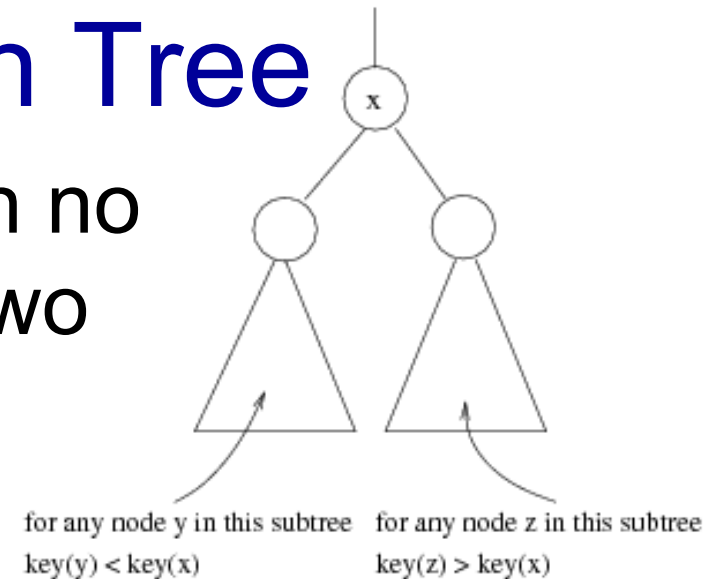
Tree



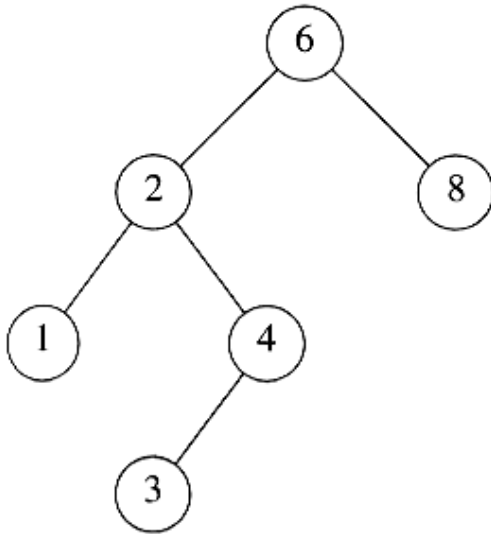
- **Child and parent**
 - Every node except the root has one parent
 - A node can have an arbitrary number of children
- **Leaves:** Nodes with no children
- **Sibling:** nodes with same parent
- **Path Length:** number of edges on the path
- **Depth of a node:** length of the unique path from the root to that node. The depth of a tree is equal to the depth of the deepest leaf
- **Height of a node:** length of the longest path from that node to a leaf, all leaves are at height 0, The height of a tree is equal to the height of the root
- **Ancestor and descendant**

Binary Search Tree


- **Binary tree:** A tree in which no node can have more than two children
- **Binary search tree:**
 - Stores keys in the nodes in a way so that searching, insertion and deletion can be done efficiently.
 - For every node X , all the keys in its left subtree are smaller than the key value in X , and all the keys in its right subtree are larger than the key value in X



Binary Search Tree

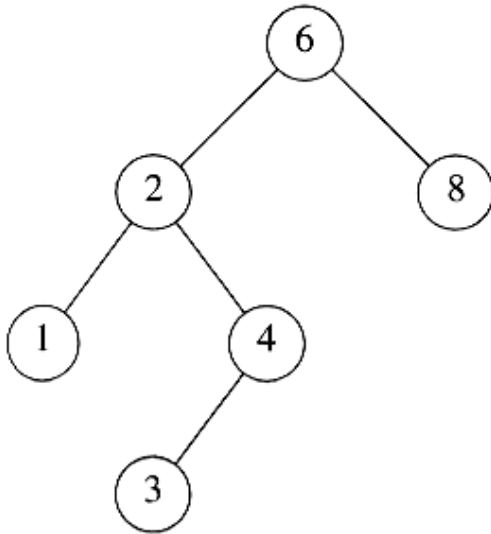


A binary search tree




Not a binary search tree
WHY?

Binary Search Tree



A binary search tree



Not a binary search tree
Because $7 > 6$

Binary Search Tree Operations

- **Tree Traversal:** Used to print out the data in a tree in a certain order
- **Pre-order traversal (node-left-right)**
 - Print the data at the root
 - Recursively print out all data in the left subtree
 - Recursively print out all data in the right subtree
- See also, **post-order: left-right-node** and **in-order:left-node-right** traversal.

Binary-search-tree sort

$T \leftarrow \emptyset$ ▷ Create an empty BST

for $i = 1$ to n

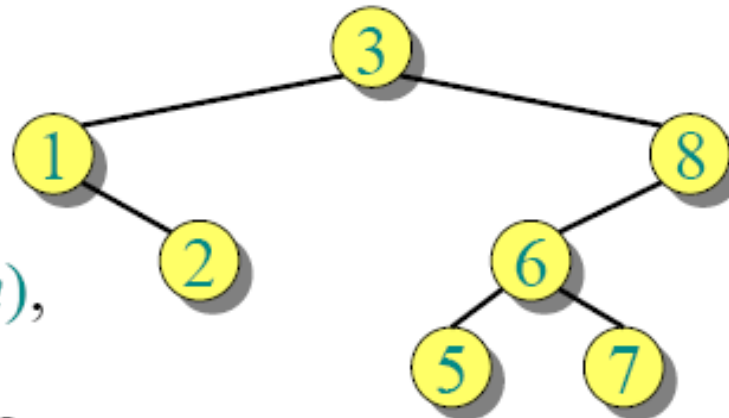
do TREE-INSERT ($T, A[i]$)

Perform an inorder tree walk (traversal) of T .

Example:

$A = [3 \ 1 \ 8 \ 2 \ 6 \ 7 \ 5]$

Tree-walk time = $O(n)$,
but how long does it
take to build the BST?



Node depth

The depth of a node = the number of comparisons made during TREE-INSERT. Assuming all input permutations are equally likely, we have

Average node depth

$$\begin{aligned} &= \frac{1}{n} E \left[\sum_{i=1}^n (\# \text{ comparisons to insert node } i) \right] \\ &= \frac{1}{n} O(n \lg n) \quad (\text{quicksort analysis}) \\ &= O(\lg n) . \end{aligned}$$

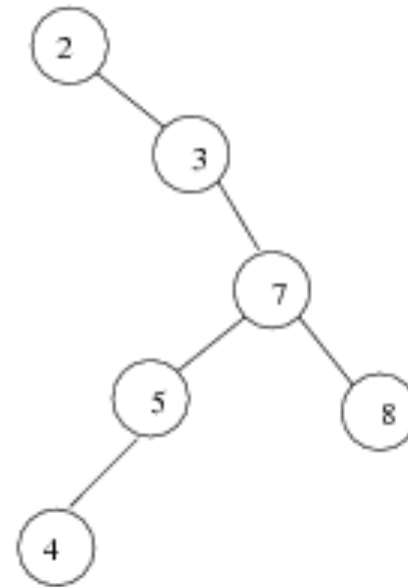
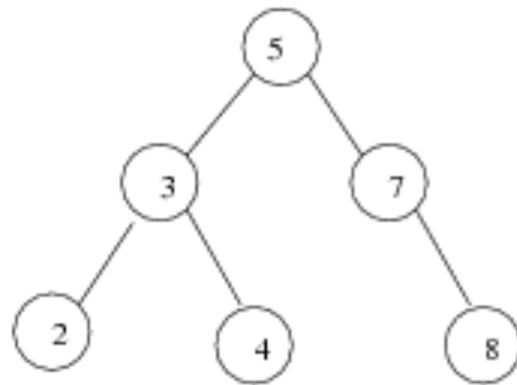
Height of a randomly built binary search tree

Outline of the analysis:

- Use *Jensen's inequality*, which says that $f(E[X]) \leq E[f(X)]$ for any convex function f and random variable X .
- Analyze the *exponential height* of a randomly built BST on n nodes, which is the random variable $Y_n = 2^{X_n}$, where X_n is the random variable denoting the height of the BST.
- Prove that $2^{E[X_n]} \leq E[2^{X_n}] = E[Y_n] = O(n^3)$, and hence that $E[X_n] = O(\lg n)$.

Binary Search Tree Operations

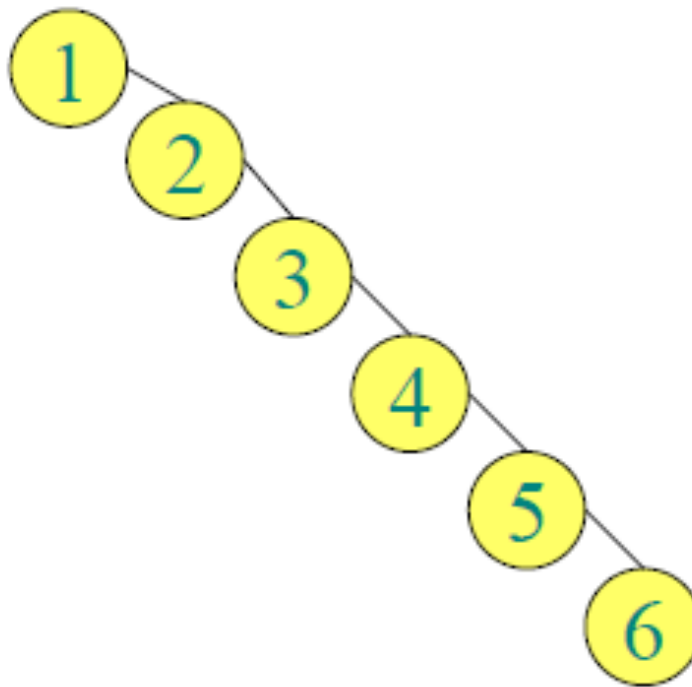
- Linear access time of linked lists is prohibitive
- Does there exist any simple data structure for which the running time of most operations (search, insert, delete) is $O(\log N)$?



Two binary search trees representing the same set.

Average depth of a node is $O(\log N)$;
maximum depth of a node is $O(N)$

Balanced search trees, or how to avoid **this** even in the worst case



Balanced search tree

Balanced search tree: A search-tree data structure for which a height of $O(\lg n)$ is guaranteed when implementing a dynamic set of n items.

Examples:

- AVL trees
- 2-3 trees
- 2-3-4 trees
- B-trees
- Red-black trees

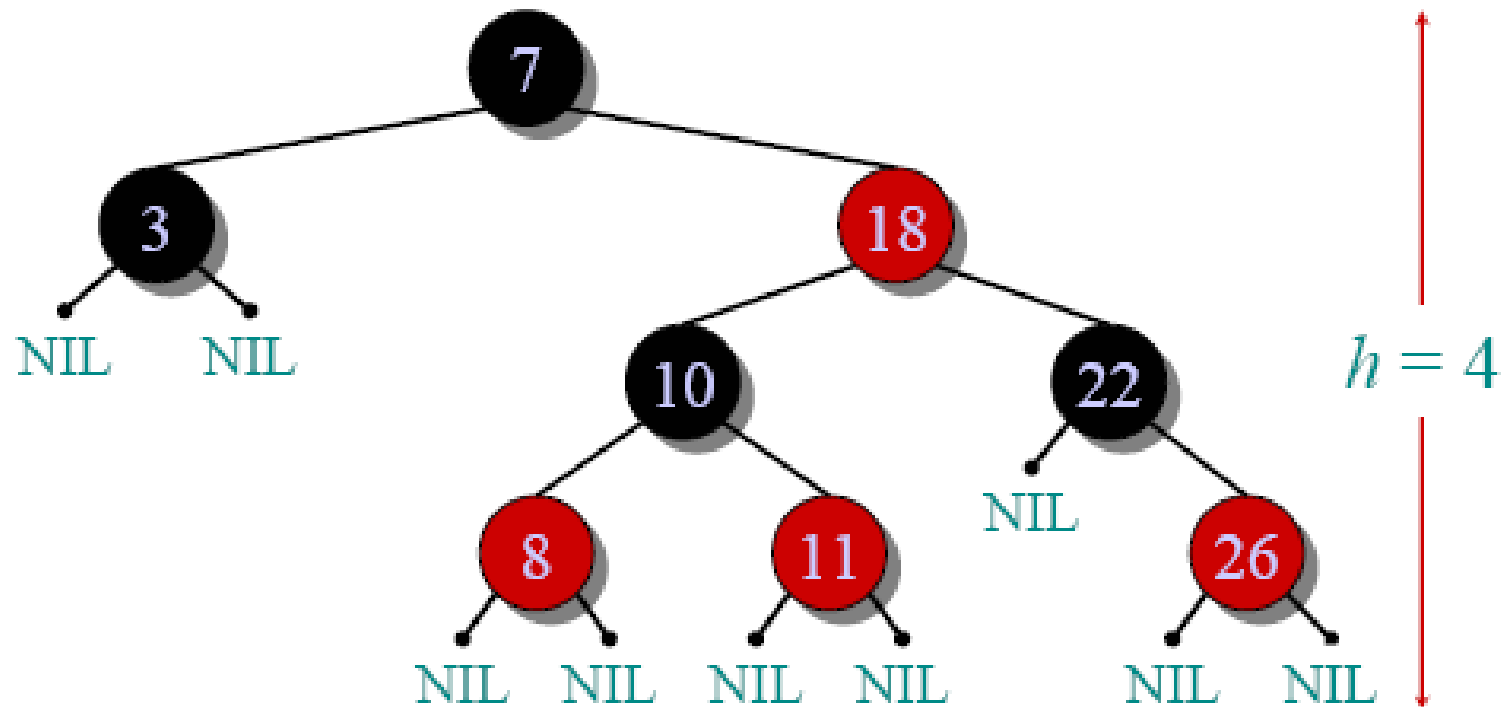
Red-black trees

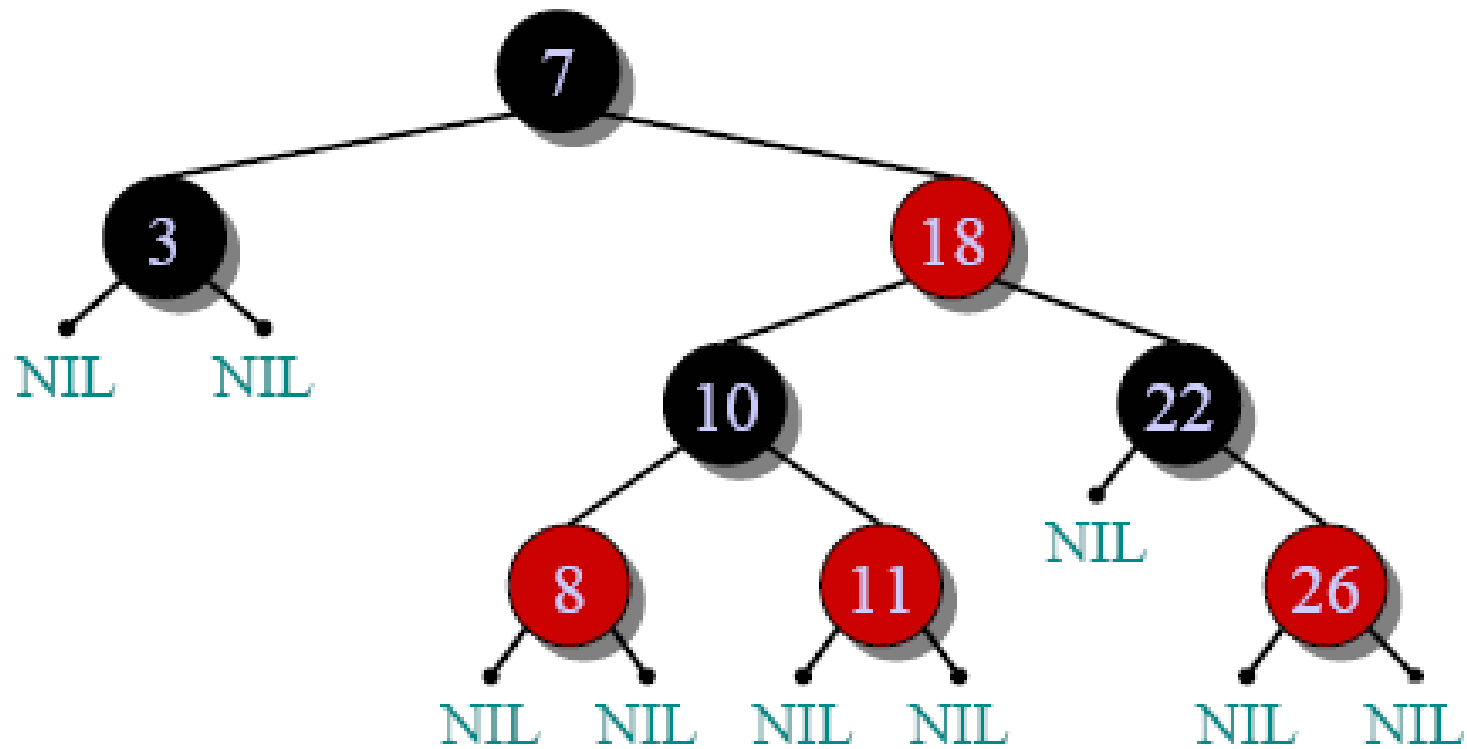
BSTs (Binary Search Tree) with an extra one-bit color field in each node.

Red-black properties:

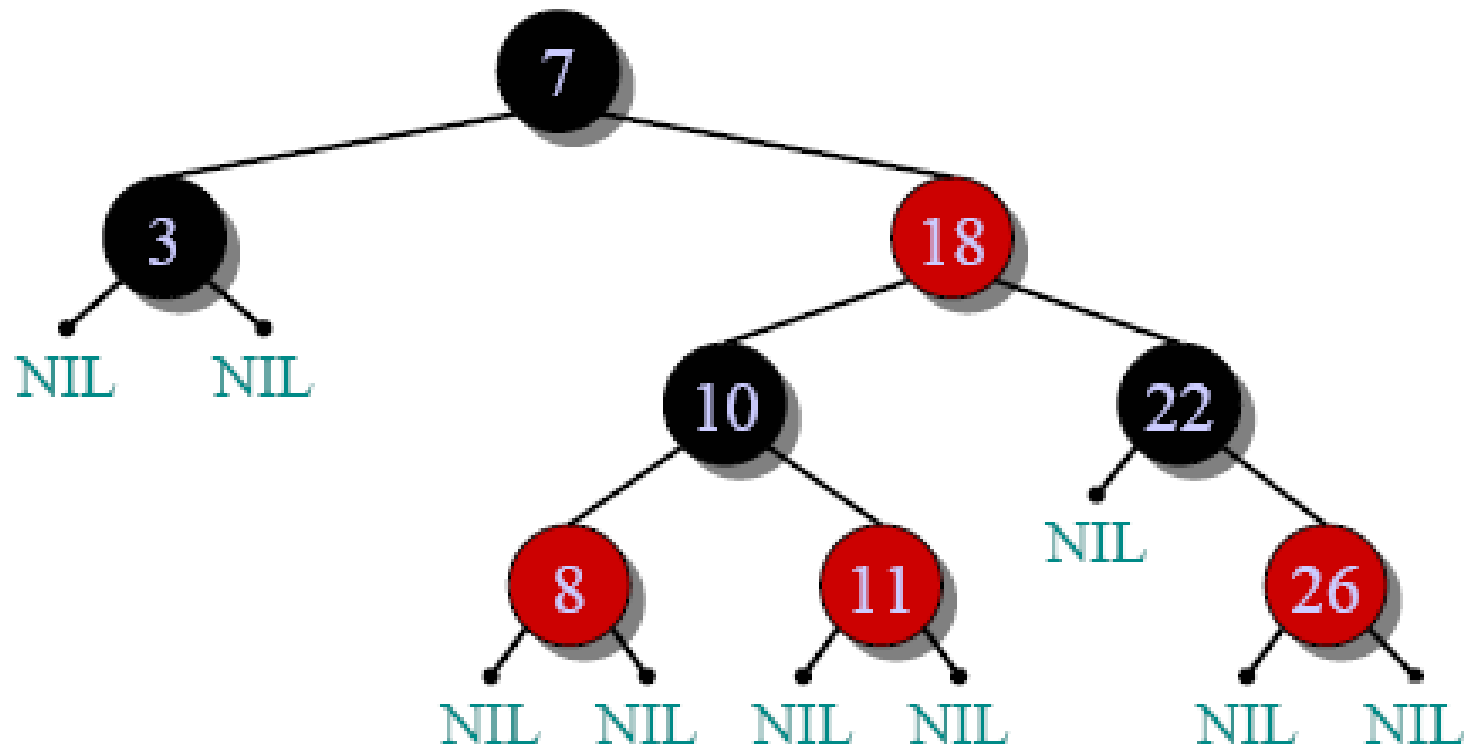
1. Every node is either red or black.
2. The root and leaves (NIL's) are black.
3. If a node is red, then its parent is black.
4. All simple paths from any node x to a descendant leaf have the same number of black nodes = $\text{black-height}(x)$.

Red-black tree example






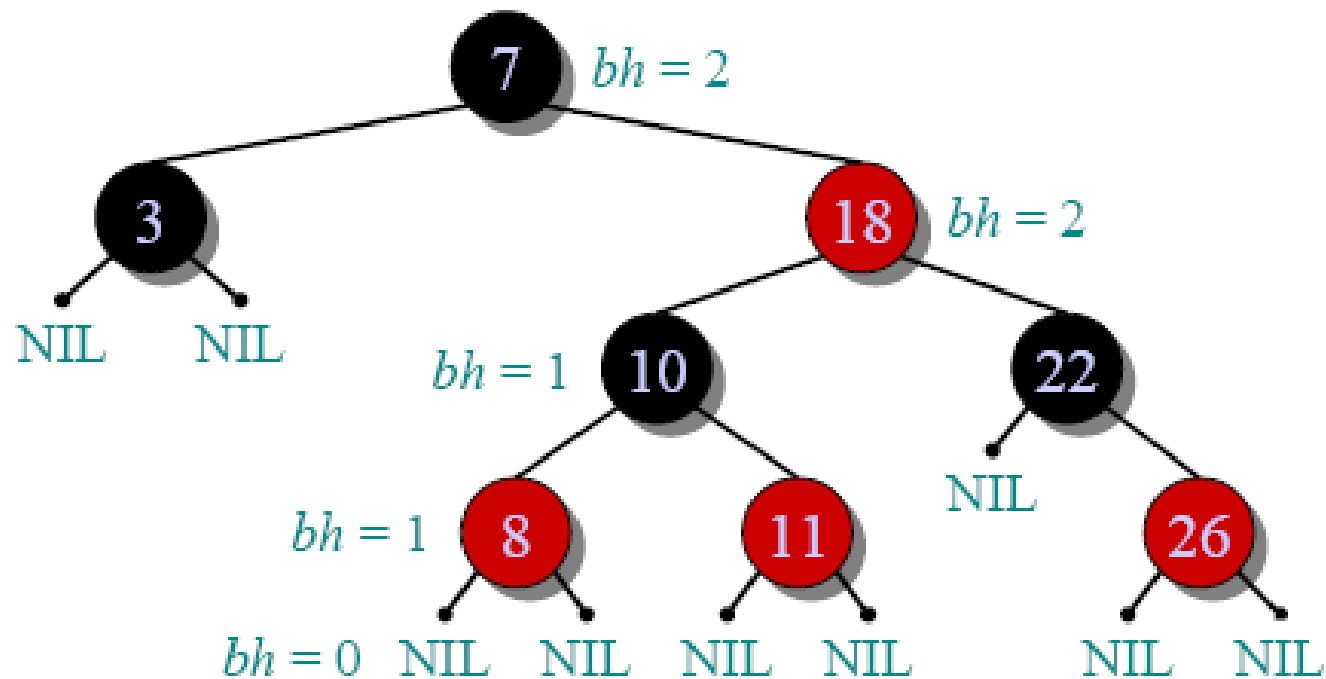
1. Every node is either red or black.



2. The root and leaves (NIL's) are black.



3. If a node is red, then its parent is black.




4. All simple paths from any node x to a descendant leaf have the same number of black nodes = *black-height*(x).

What properties would we like to prove about red-black trees ?

They always have **$O(\log n)$ height**

There is an **$O(\log n)$ time insertion** procedure which preserves the red-black properties

- Is it true that, after we add a new element to a tree, we can always recolor the tree to keep it red-black ?

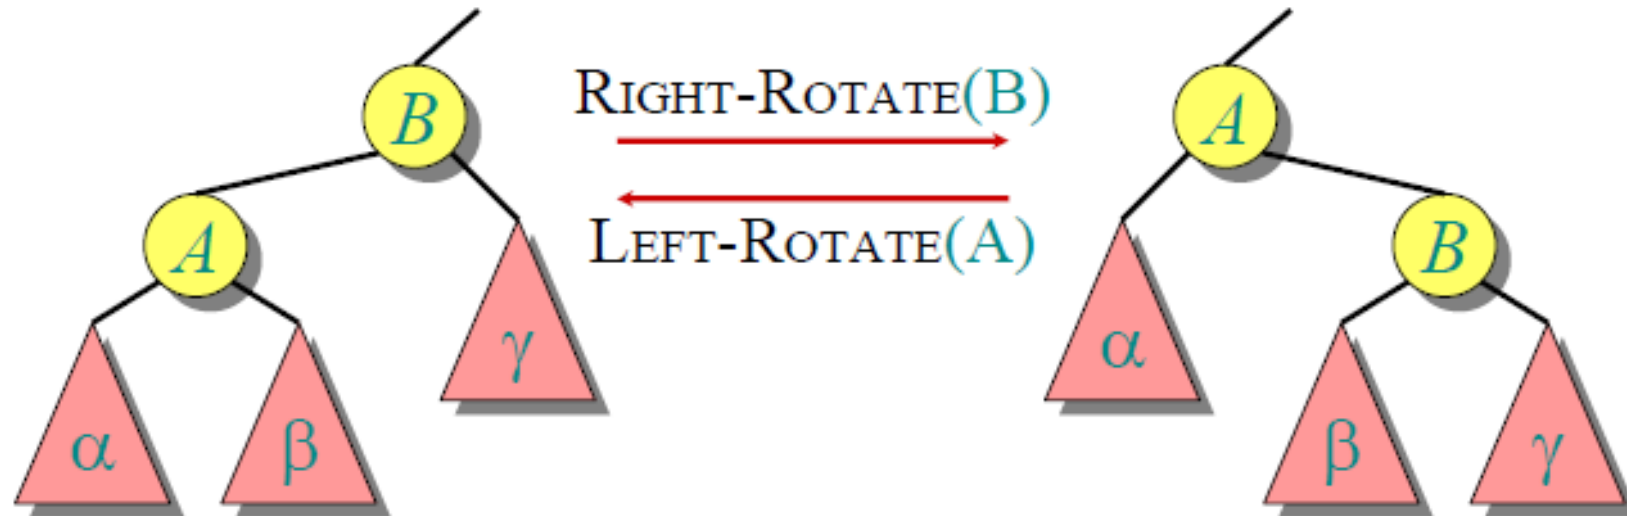


- Is it true that, after we add a new element to a tree, we can always recolor the tree to keep it red-black ?

NO. After insertions, sometimes we need to juggle nodes around



Rotations





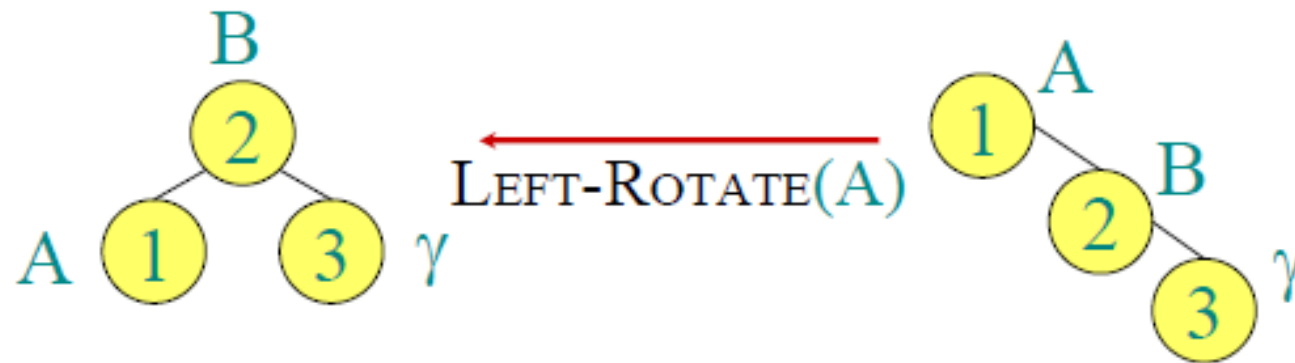
Rotations maintain the inorder ordering of keys:

- $a \in \alpha, b \in \beta, c \in \gamma \Rightarrow a \leq A \leq b \leq B \leq c.$

A rotation can be performed in $O(1)$ time.



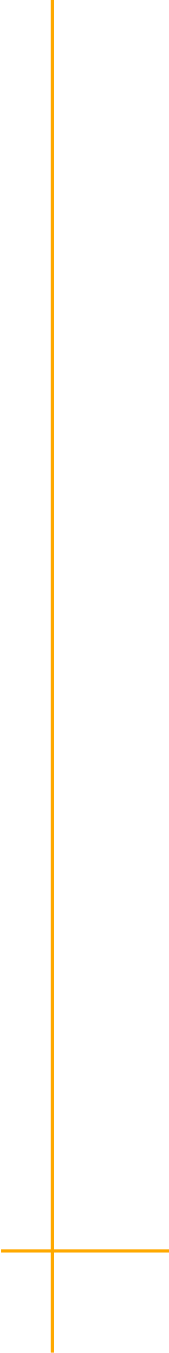
Rotations can reduce height





Red-black tree wrap-up

- Can show how
 - $O(\log n)$ re-colorings
 - 1 rotationcan restore red-black properties after an insertion
- Instead, we will see 2-3 trees (but will come back to red-black trees at the end)




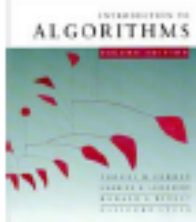
2-3 and 2-3-4 Trees



2-3 Trees

- The simplest balanced trees on the planet!
- Although a little bit more wasteful
- Degree of each node is either 2 or 3
- Keys are in the leaves
- All leaves have equal depth
- Leaves are sorted
- Each node x contains maximum key in the sub-tree, denoted $x.max$





Internal nodes

- Internal nodes:
 - Values:
 - `x.max`: maximum key in the sub-tree
 - Pointers:
 - `left[x]`
 - `mid[x]`
 - `right[x]` : can be null
 - `p[x]` : can be null for the root
 - ...
- Leaves:
 - `x.max` : the key



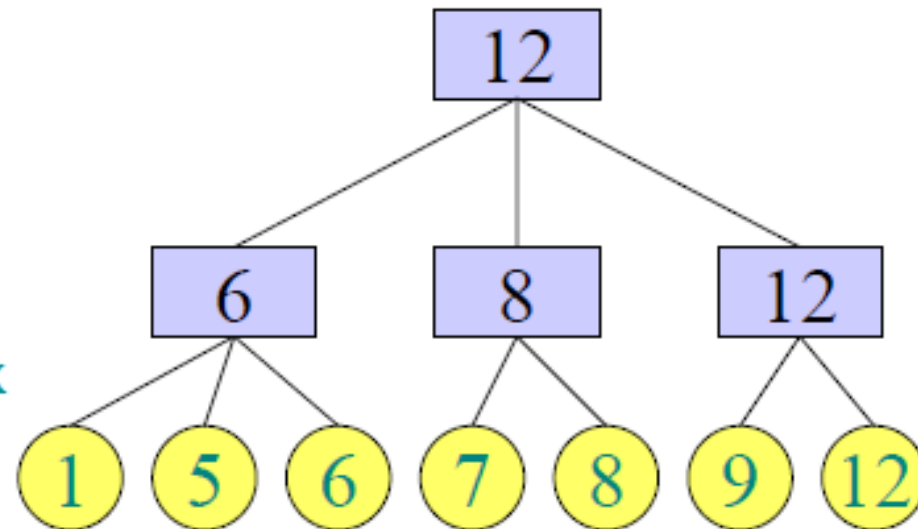
Height of 2-3 tree

- What is the maximum height h of a 2-3 tree with n nodes ?
- Alternatively, what is the minimum number of nodes in a 2-3 tree of height h ?
- It is $1+2+2^2+2^3+\dots+2^h = 2^{h+1}-1$
- $n \geq 2^{h+1}-1 \Rightarrow h = O(\log n)$
- Full binary tree is the worst-case example!



Searching[^]

- How can we search for a key k ?
- Search(x, k):
- If $x = \text{NIL}$ then return NIL
 - Else if x is a leaf then
 - If $x.\text{max} = k$ then return x
 - Else return NIL
 - Else
 - If $k \leq \text{left}[x].\text{max}$ then Search($\text{left}[x], k$)
 - Else if $k \leq \text{mid}[x].\text{max}$ then Search($\text{mid}[x], k$)
 - Else Search($\text{right}[x], k$)

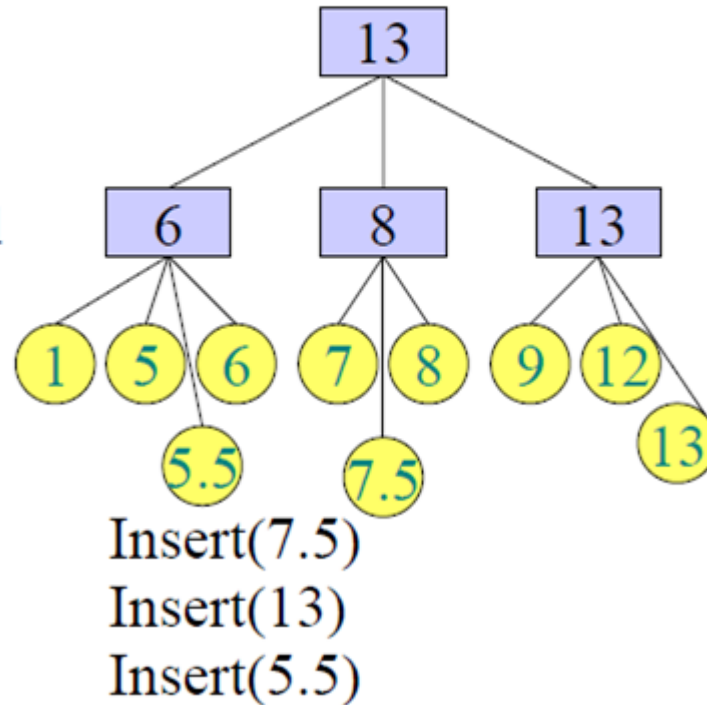


Search(8)
Search(13)

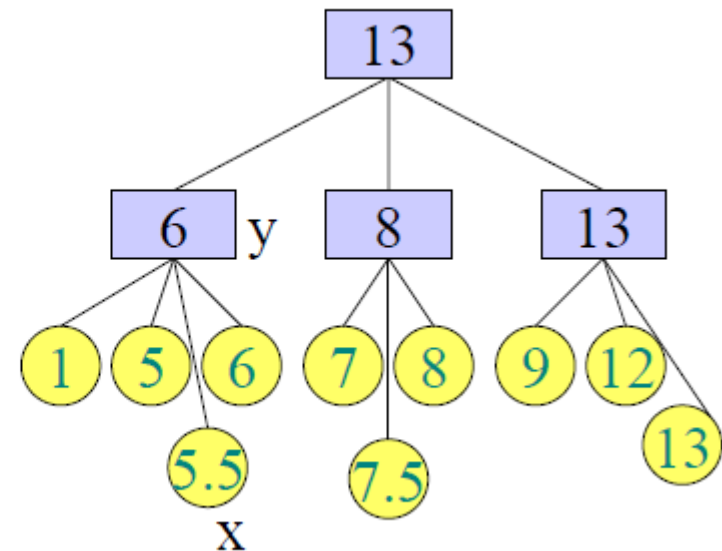


Insertion

- How to insert x ?
- Perform Search for the key of x
- Let y be the last internal node
- Insert x into y in a sorted order
- At the end, update the max values on the path to root



- If y has 4 children, then Split(y)

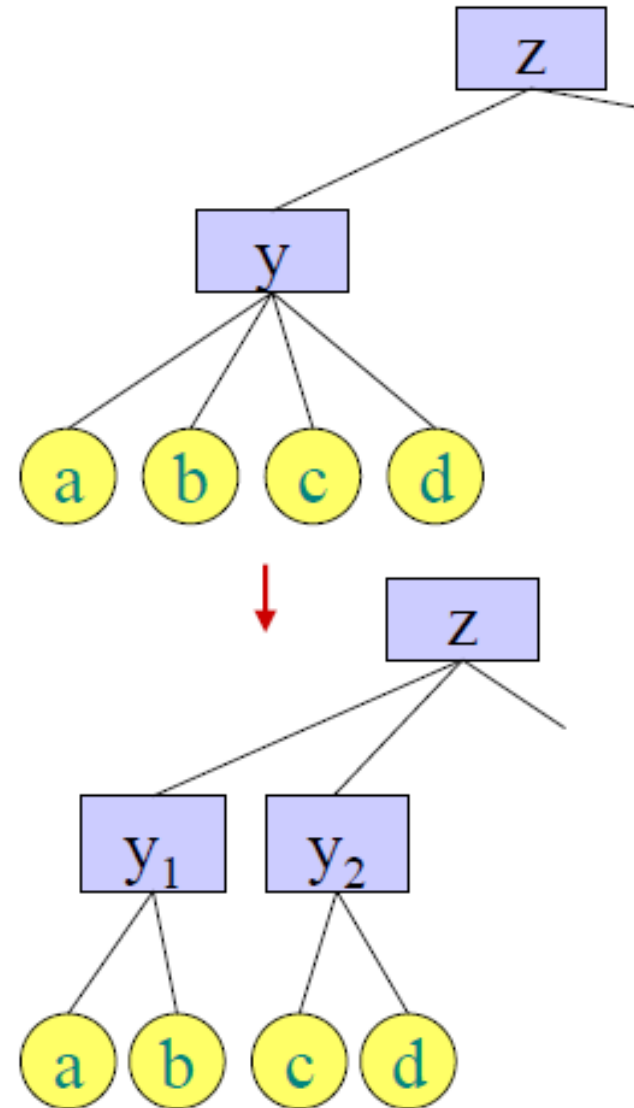


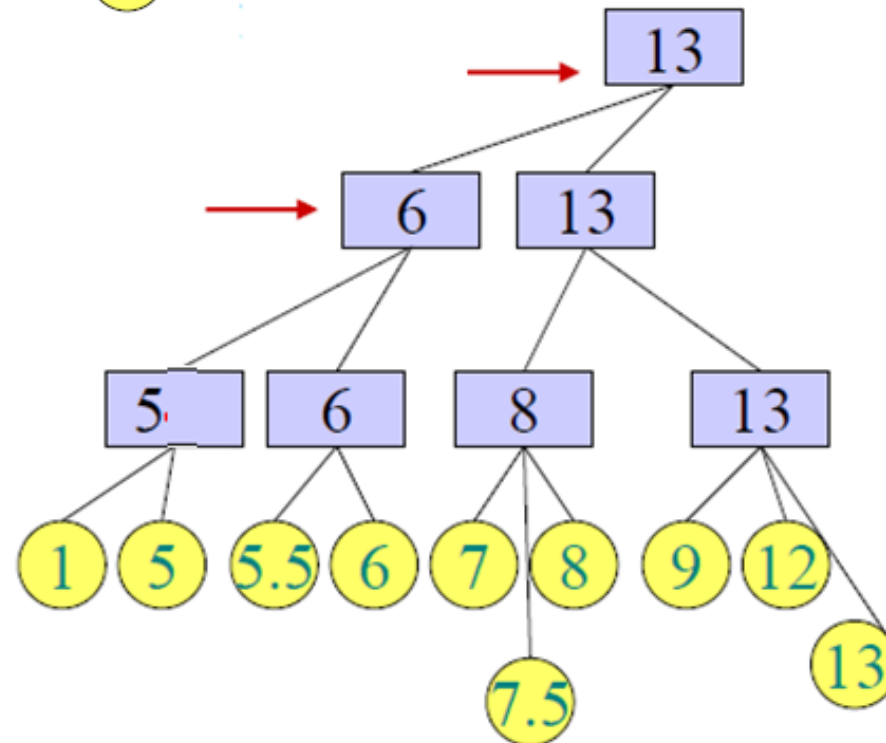
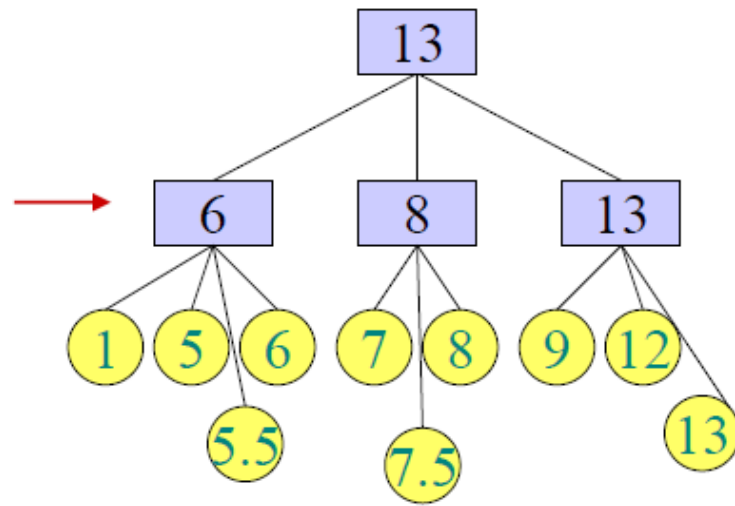


Split

- Split y into two nodes y_1, y_2
- Both are linked to $z = \text{parent}(y)^*$
- If z has 4 children, split z

*If y is a root, then create new $\text{parent}(y) = \text{new root}$



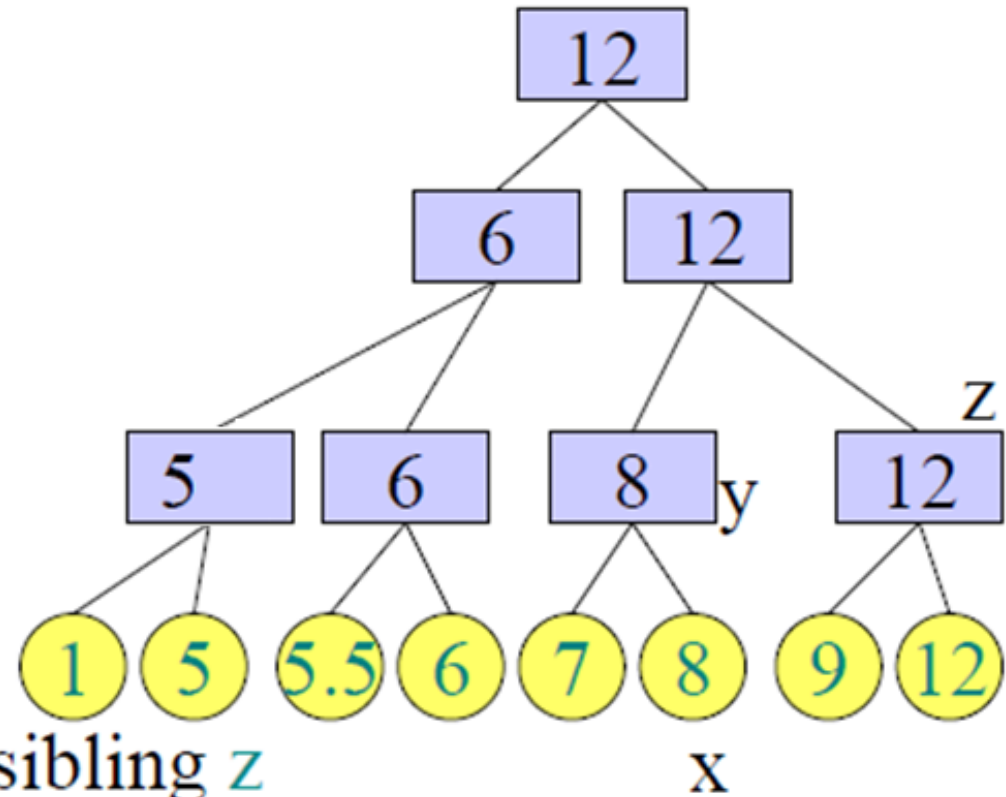


- Insert and Split preserve heights, unless new root is created, in which case all heights are increased by 1
- After Split, all nodes have 2 or 3 children
- Everything takes $O(\log n)$ time



Delete

- How to delete x ?
- Let $y = p(x)$
- Remove x from y
- If y has 1 child:
 - Remove y
 - Attach to y 's sibling z

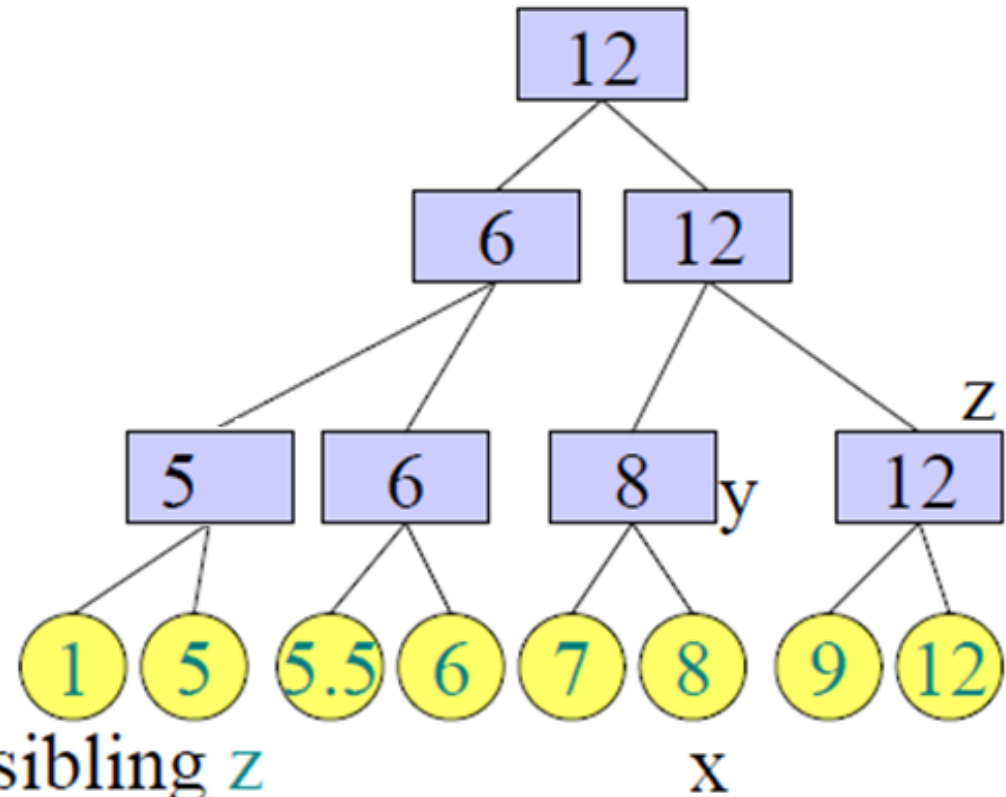


Delete(8)



Delete

- How to delete x ?
- Let $y = p(x)$
- Remove x from y
- If y has 1 child:
 - Remove y
 - Attach to y 's sibling z



Delete(8)

If z has 4 children, then Split(z)

INCOMPLETE – SEE THE END FOR FULL VERSION




Summing up

- 2-3 Trees:
 - $O(\log n)$ depth \Rightarrow Search in $O(\log n)$ time
 - Insert, Delete (and Split) in $O(\log n)$ time
- We will now see 2-3-4 trees
 - Same idea, but:
 - Each parent has 2,3 or 4 children
 - Keys in the inner nodes
 - More complicated procedures



2-3-4 Trees





Back to Red and Black Trees

Height of a red-black tree


Theorem. A red-black tree with n keys has height

$$h \leq 2 \lg(n + 1).$$

Proof. (The book uses induction. Read carefully.)

INTUITION:

Merge red nodes into their black parents.



Height of a red-black tree


Theorem. A red-black tree with n keys has height

$$h \leq 2 \lg(n + 1).$$

Proof. (The book uses induction. Read carefully.)

INTUITION:

Merge red nodes into their black parents.



Height of a red-black tree


Theorem. A red-black tree with n keys has height

$$h \leq 2 \lg(n + 1).$$

Proof. (The book uses induction. Read carefully.)

INTUITION:

Merge red nodes into their black parents.



Height of a red-black tree

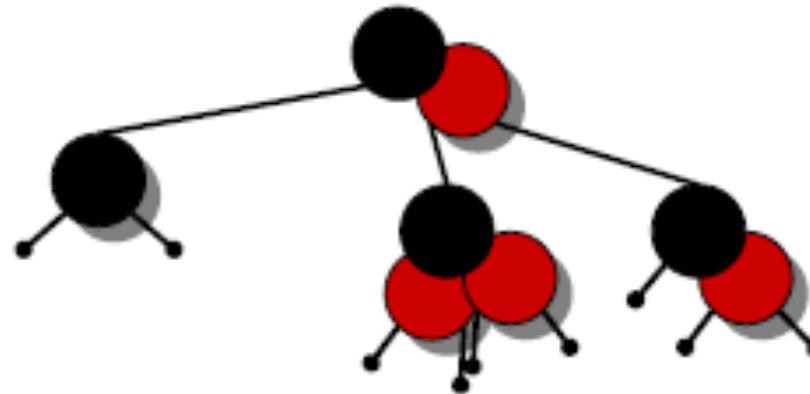
Theorem. A red-black tree with n keys has height

$$h \leq 2 \lg(n + 1).$$

Proof. (The book uses induction. Read carefully.)

INTUITION:

Merge red nodes into their black parents.



Height of a red-black tree

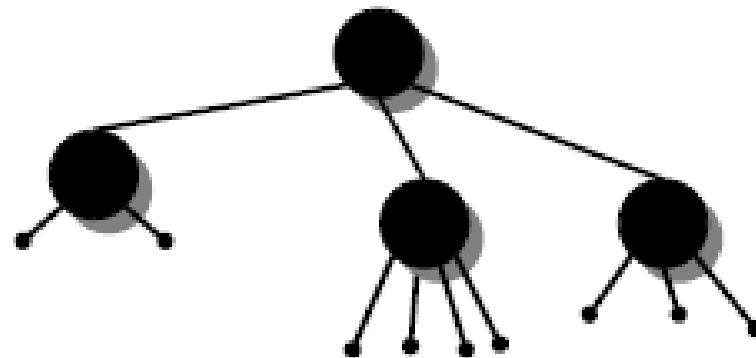
Theorem. A red-black tree with n keys has height

$$h \leq 2 \lg(n + 1).$$

Proof. (The book uses induction. Read carefully.)

INTUITION:

Merge red nodes into their black parents.



Height of a red-black tree


Theorem. A red-black tree with n keys has height

$$h \leq 2 \lg(n + 1).$$

Proof. (The book uses induction. Read carefully.)

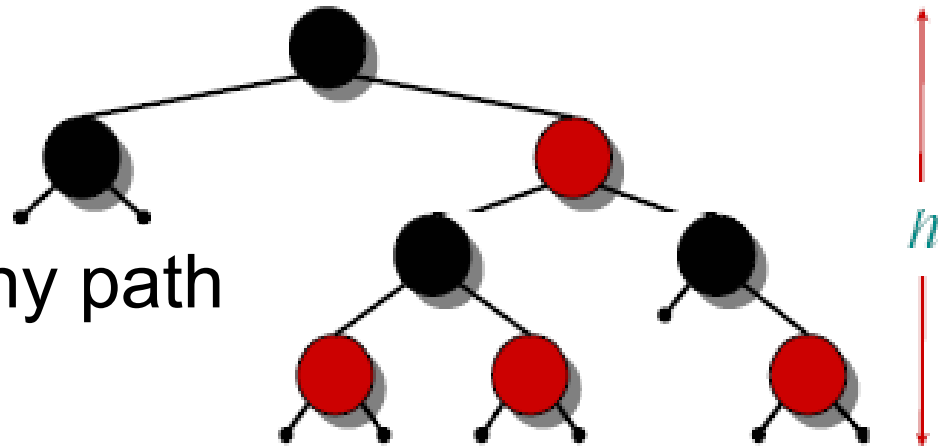
INTUITION:

- Merge red nodes into their black parents.
- This process produces a tree in which each node has 2, 3, or 4 children.
- The 2-3-4 tree has uniform depth h' of leaves.



Height of a red-black tree

- We have $h' \geq h/2$, since at most half the leaves on any path are red.

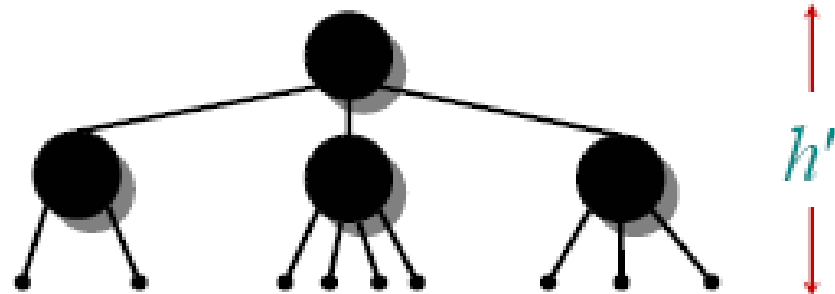


- The number of leaves in each tree is $n + 1$

$$\Rightarrow n + 1 \geq 2^{h'}$$

$$\Rightarrow \lg(n + 1) \geq h' \geq h/2$$

$$\Rightarrow h \leq 2 \lg(n + 1). \quad \square$$



Query operations

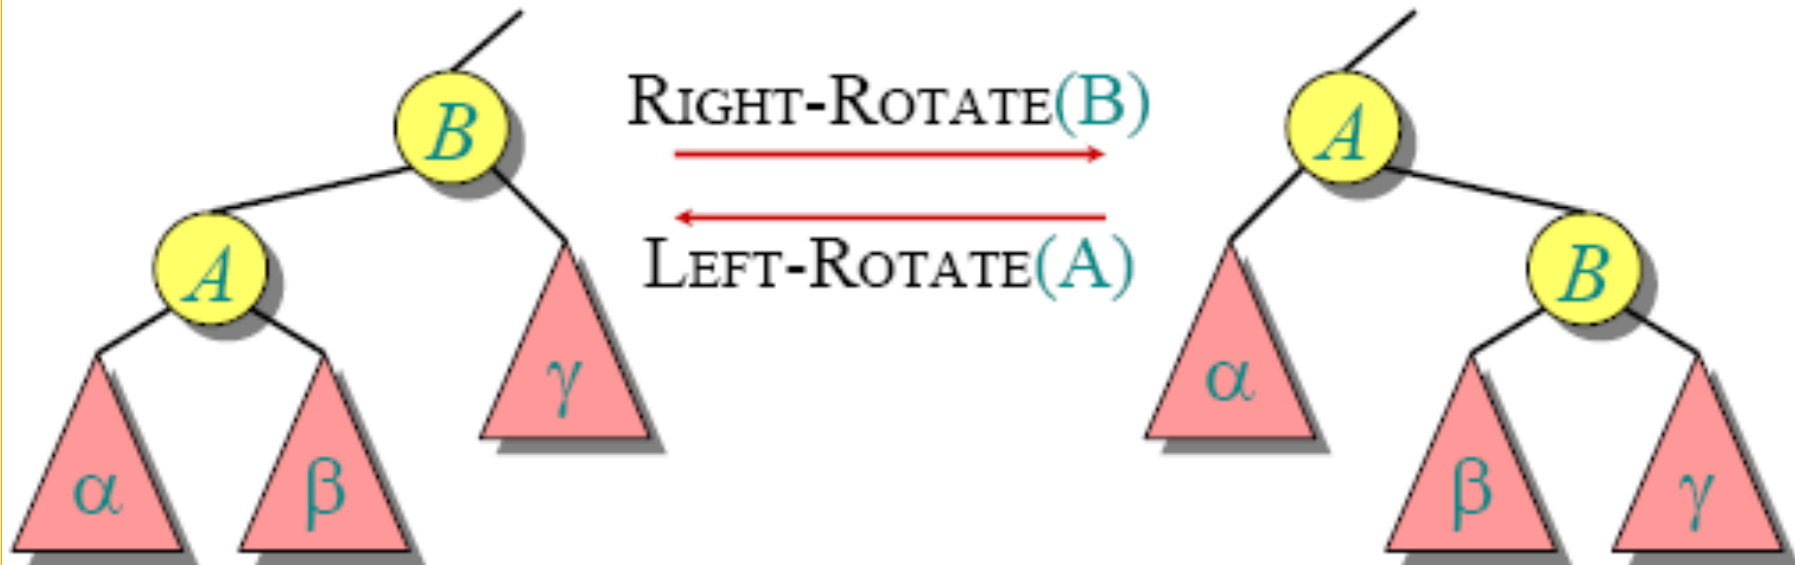
Corollary. The queries SEARCH, MIN, MAX, SUCCESSOR, and PREDECESSOR all run in $O(\lg n)$ time on a red-black tree with n nodes.

Modifying operations

The operations INSERT and DELETE cause modifications to the red-black tree:

- the operation itself,
- color changes,
- restructuring the links of the tree via ***“rotations”***.

Rotations



Rotations maintain the inorder ordering of keys:

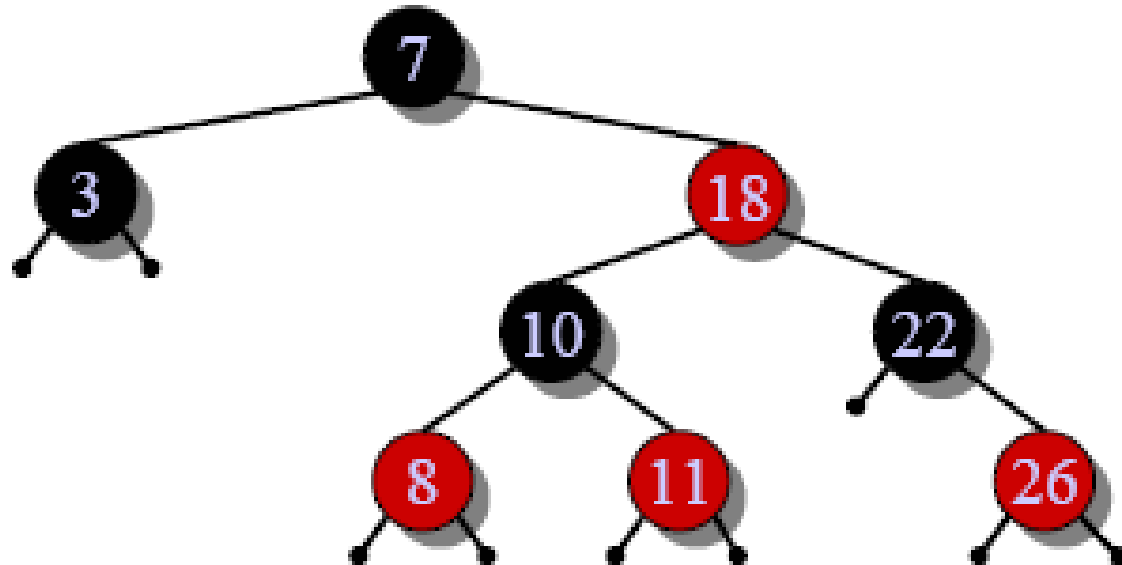
- $a \in \alpha, b \in \beta, c \in \gamma \Rightarrow a \leq A \leq b \leq B \leq c.$

A rotation can be performed in $O(1)$ time.

Insertion into a red-black tree

IDEA: Insert x in tree. Color x red. Only redblack property 3 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

Example:

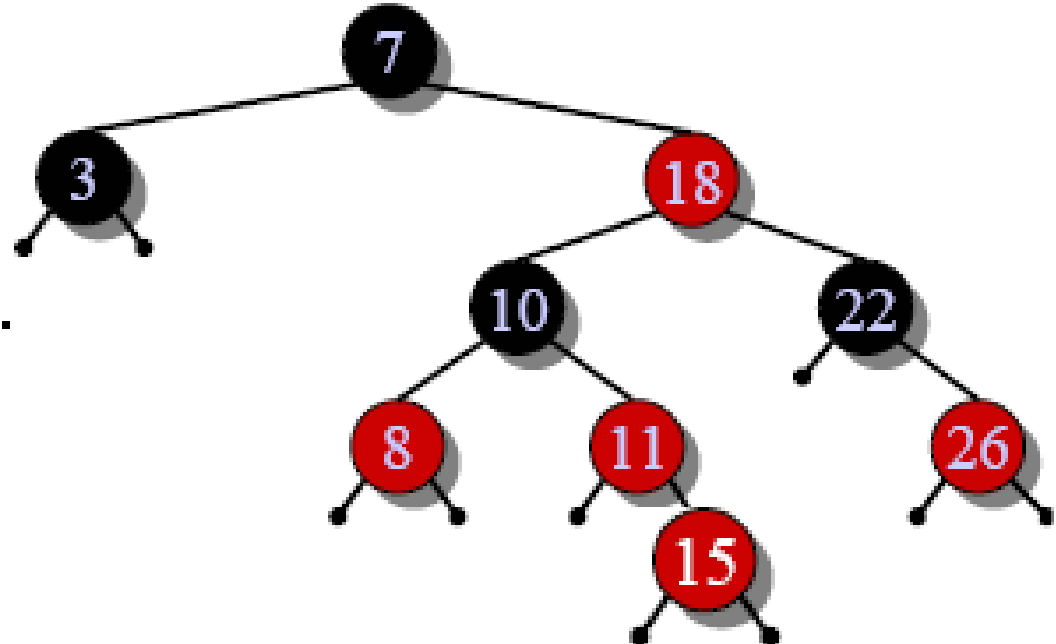


Insertion into a red-black tree

IDEA: Insert x in tree. Color x red. Only redblack property 3 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

Example:

- Insert $x = 15$.
- Recolor, moving the violation up the tree.

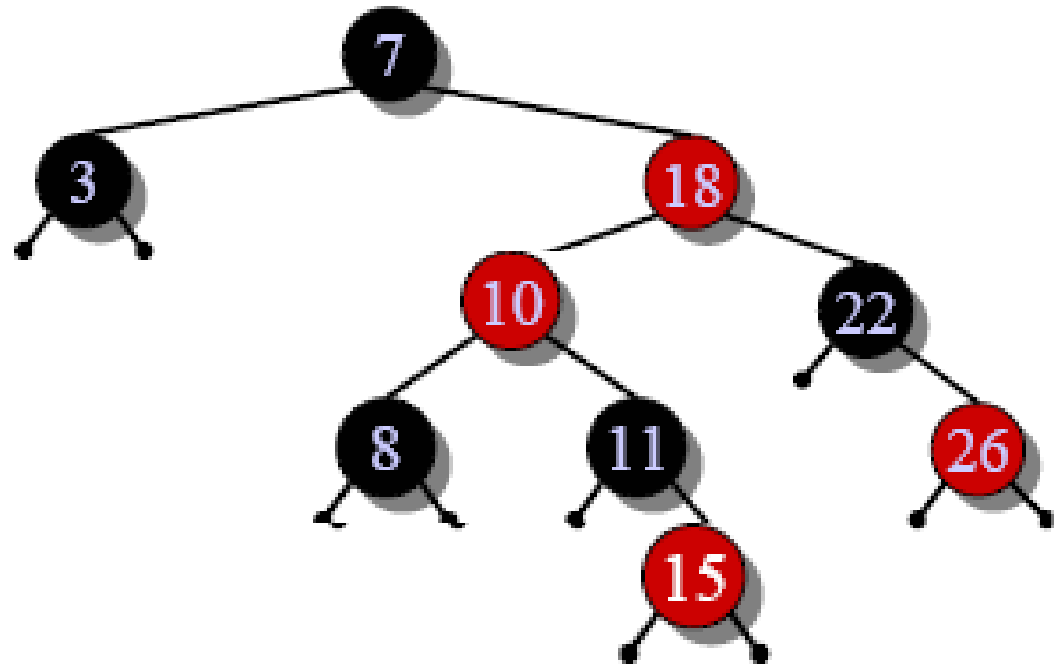


Insertion into a red-black tree

IDEA: Insert x in tree. Color x red. Only redblack property 3 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

Example:

- Insert $x = 15$.
- Recolor, moving the violation up the tree
- RIGHT-ROTATE(18)

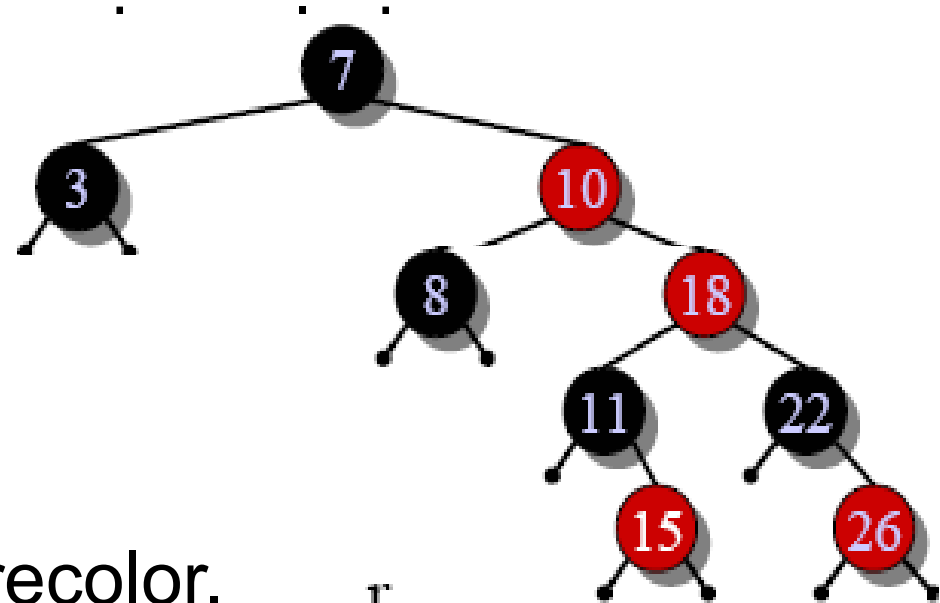


Insertion into a red-black tree

IDEA: Insert x in tree. Color x red. Only redblack property 3 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations

Example:

- Insert $x = 15$.
- Recolor, moving the violation up the tree.
- RIGHT-ROTATE(18).
- LEFT-ROTATE(7) and recolor.

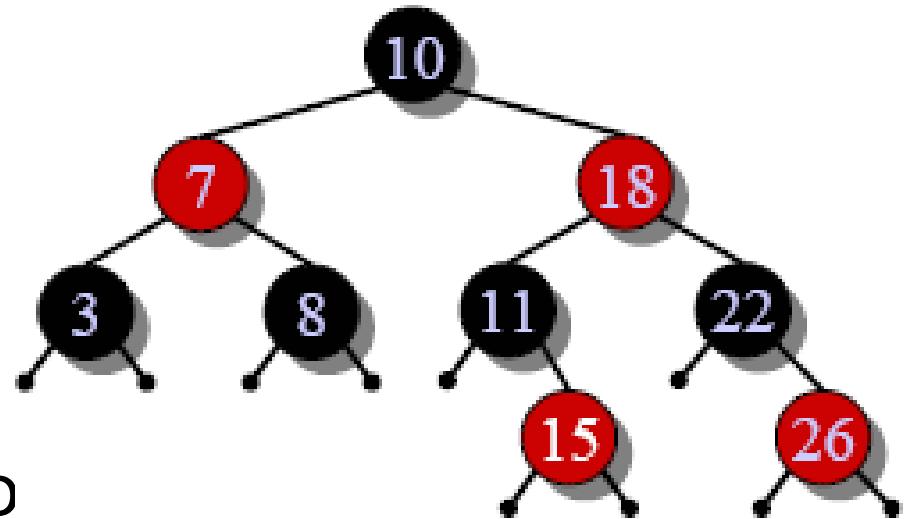


Insertion into a red-black tree

IDEA: Insert x in tree. Color x red. Only redblack property 3 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

Example:

- Insert $x = 15$.
- Recolor, moving the violation up the tree.
- RIGHT-ROTATE(18).
- LEFT-ROTATE(7) and recolor



Pseudocode

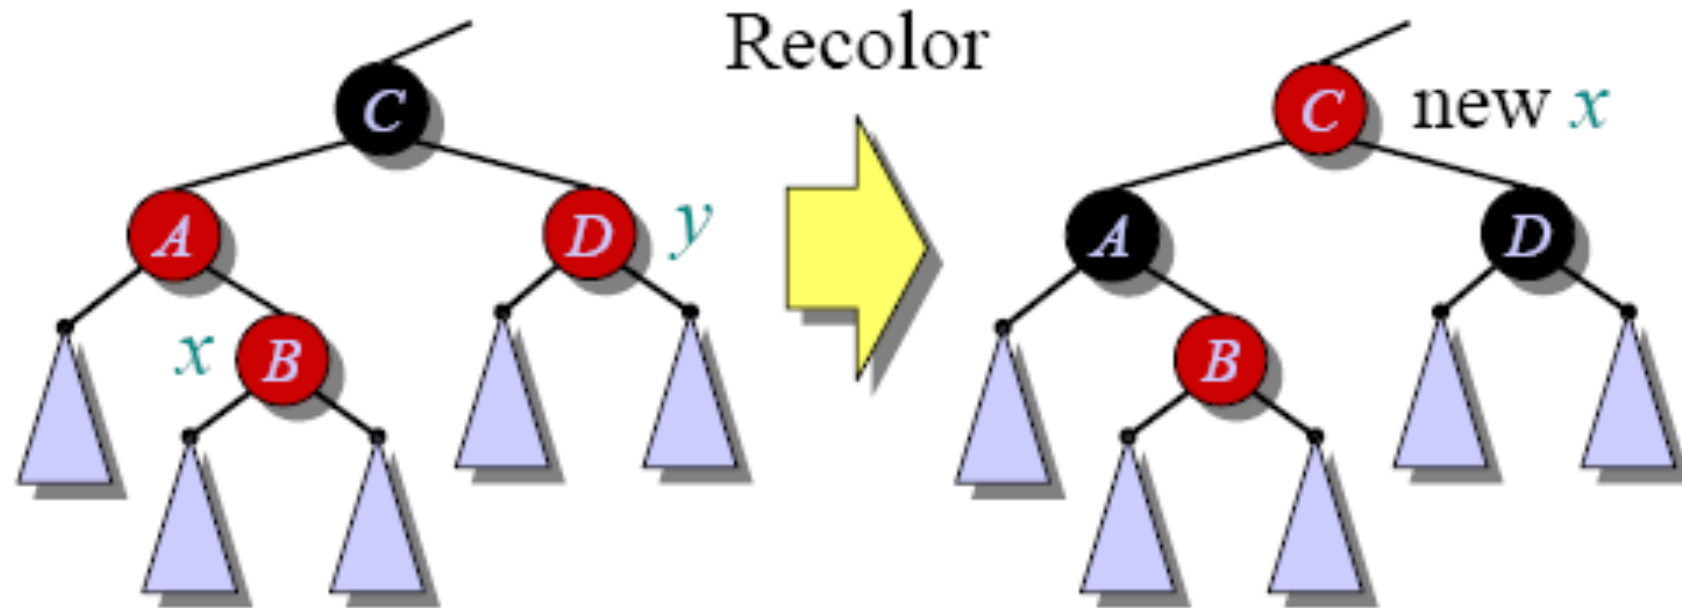
```
RB-INSERT( $T, x$ )
  TREE-INSERT( $T, x$ )
   $color[x] \leftarrow RED$   $\triangleright$  only RB property 3 can be violated
  while  $x \neq root[T]$  and  $color[p[x]] = RED$ 
    do if  $p[x] = left[p[p[x]]]$ 
      then  $y \leftarrow right[p[p[x]]]$   $\triangleright y = \text{aunt/uncle of } x$ 
        if  $color[y] = RED$ 
          then  $\langle \text{Case 1} \rangle$ 
          else if  $x = right[p[x]]$ 
            then  $\langle \text{Case 2} \rangle \triangleright \text{Case 2 falls into Case 3}$ 
             $\langle \text{Case 3} \rangle$ 
          else  $\langle \text{"then" clause with "left" and "right" swapped} \rangle$ 
         $color[root[T]] \leftarrow BLACK$ 
```

Graphical notation

Let  denote a subtree with a black root.

All 's have the same black-height.

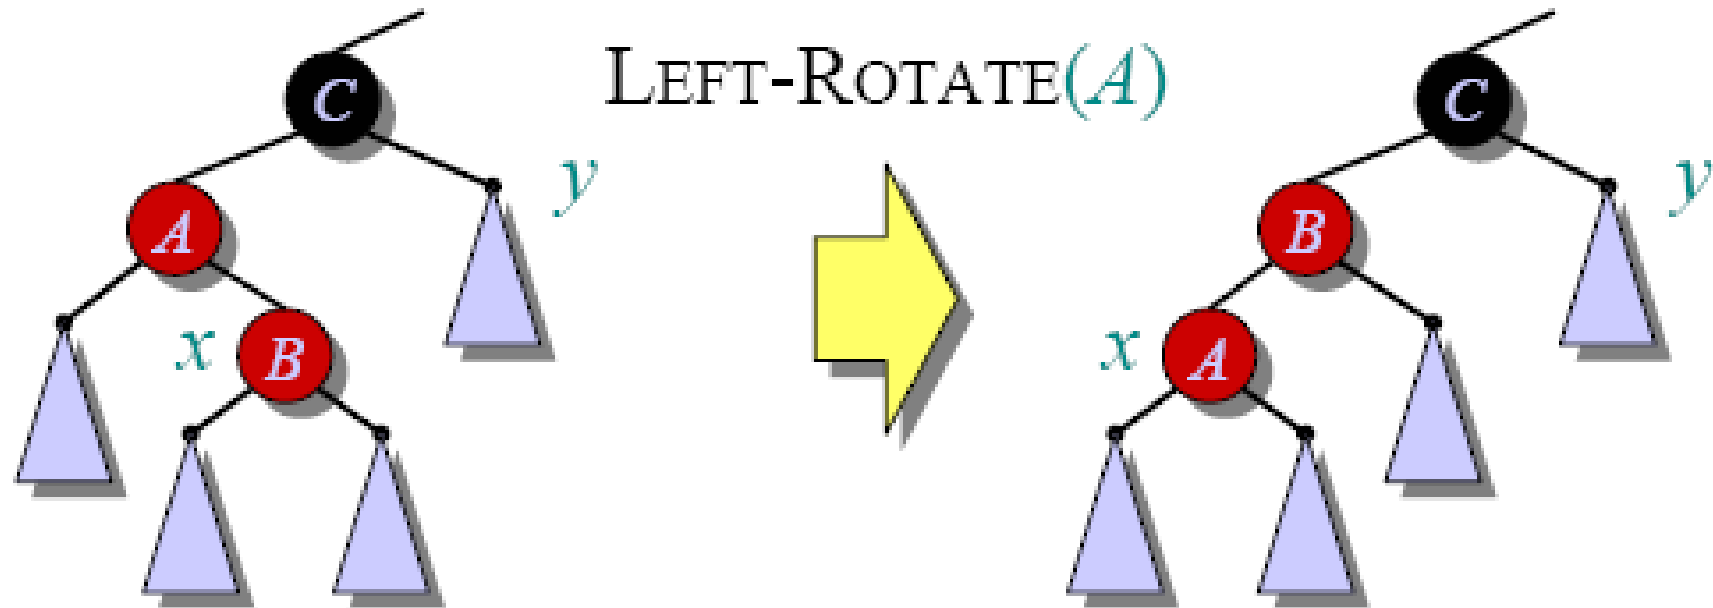
Case 1



(Or, children of A are swapped.)

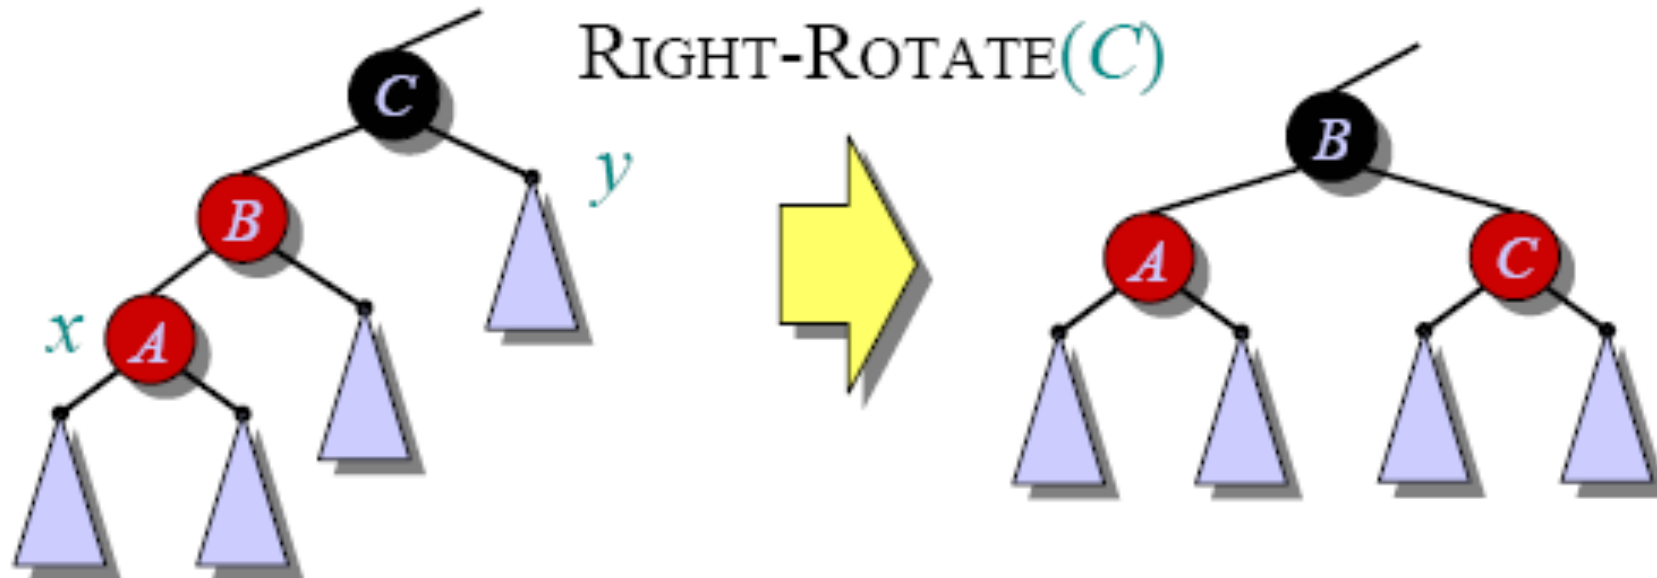
Push C 's black onto A and D , and recurse, since C 's parent may be red.

Case 2



Transform to Case 3.

Case 3



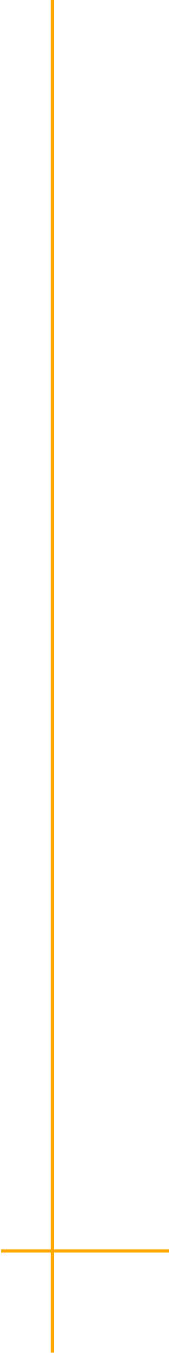
Done! No more violations of RB property 3 are possible.

Analysis

- Go up the tree performing Case 1, which only recolors nodes.
- If Case 2 or Case 3 occurs, perform 1 or 2 rotations, and terminate.

Running time: $O(\lg n)$ with $O(1)$ rotations.

RB-DELETE— same asymptotic running time and number of rotations as RB-INSERT (see textbook).

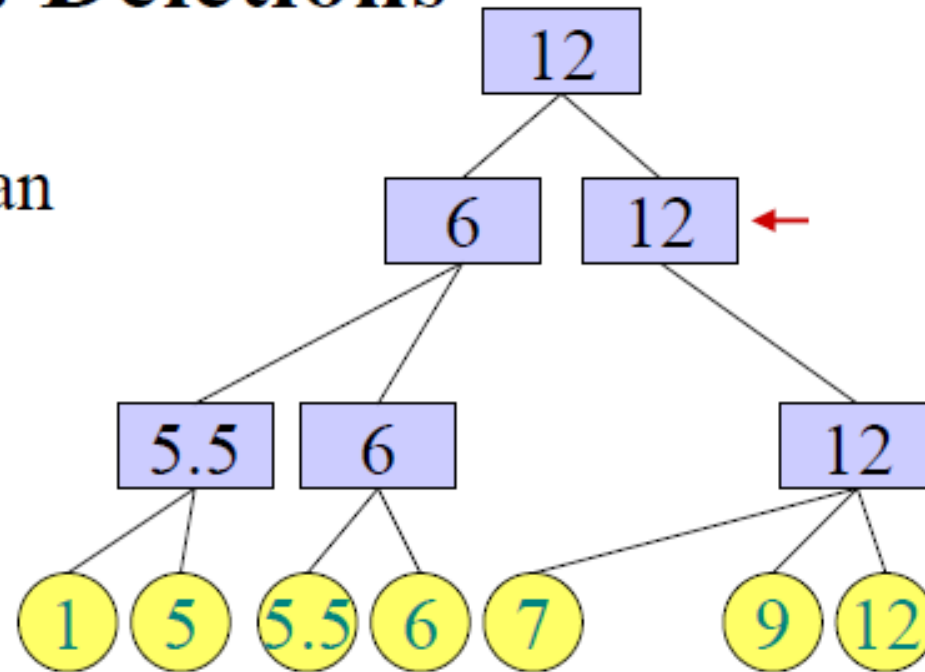


Back to 2-3 Trees



2-3 Trees: Deletions

- Problem: there is an internal node that has only 1 child



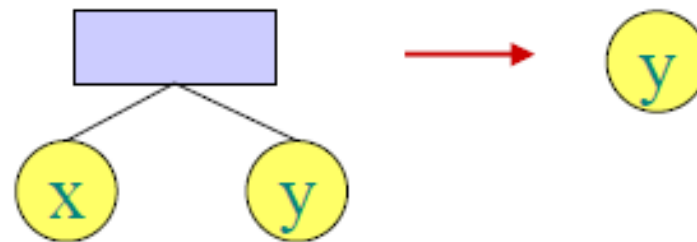


Full procedure for Delete(x)

- Special case: x is the only element in the tree: delete everything



- Not-so-special case: x is one of two elements in the tree. In this case, the procedure on the next slide will delete x

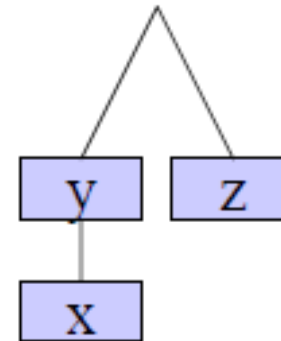


- Both NIL and y are special 2-3 trees

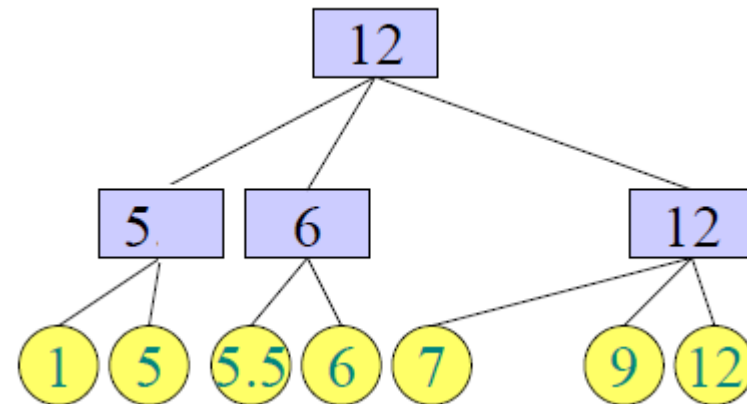
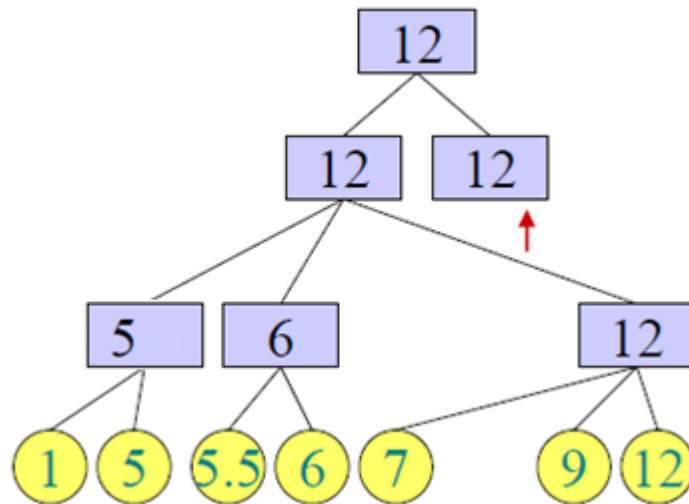
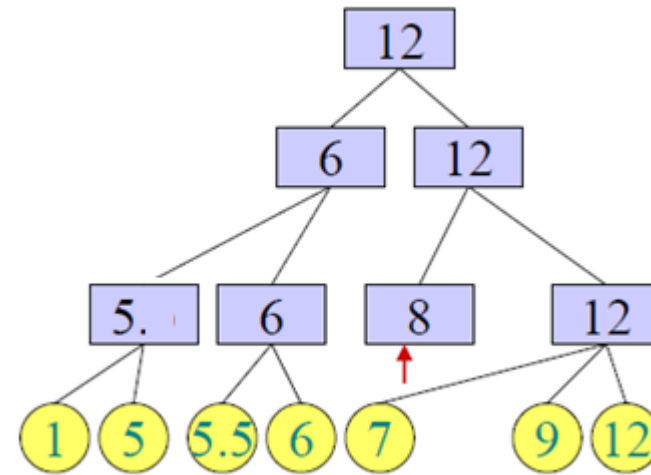
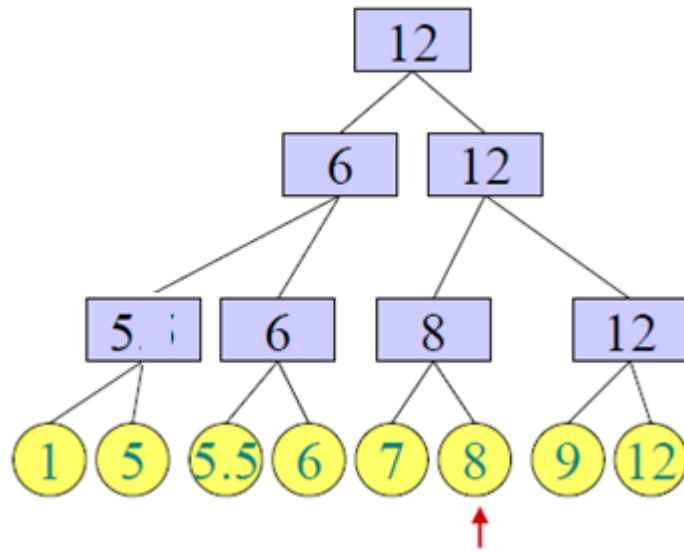


Procedure for Delete(x)

- Let $y = p(x)$
- Remove x
- If $y \neq \text{root}$ then
 - Let z be the sibling of y .
 - Assume z is the right sibling of y , otherwise the code is symmetric.
 - If y has only 1 child w left
 - Case 1: z has 3 children
 - Attach $\text{left}[z]$ as the rightmost child of y
 - Update $y.\text{max}$ and $z.\text{max}$
 - Case 2: z has 2 children:
 - Attach the child w of y as the leftmost child of z
 - Update $z.\text{max}$
 - Delete(y) (recursively*)
 - Else
 - Update max of y , $p(y)$, $p(p(y))$ and so on until root
- Else
 - If root has only one child u
 - Remove root
 - Make u the new root



*Note that the input of Delete does not have to be a leaf



Summary

Binary Search Tree (BST) review

Red and Black Trees

2-3 and 2-3-4 trees

Operations on Red and Black Trees